

# Credit Risk Probability Model for Alternative Data

**Audience:** Bati Bank risk, product, and compliance stakeholders \ **Date:** 2025-12-11 \ **Authoring repo:**

Credit-Risk-Probability-Model-for-Alternative-Data

## 1) Business Objective and Framing

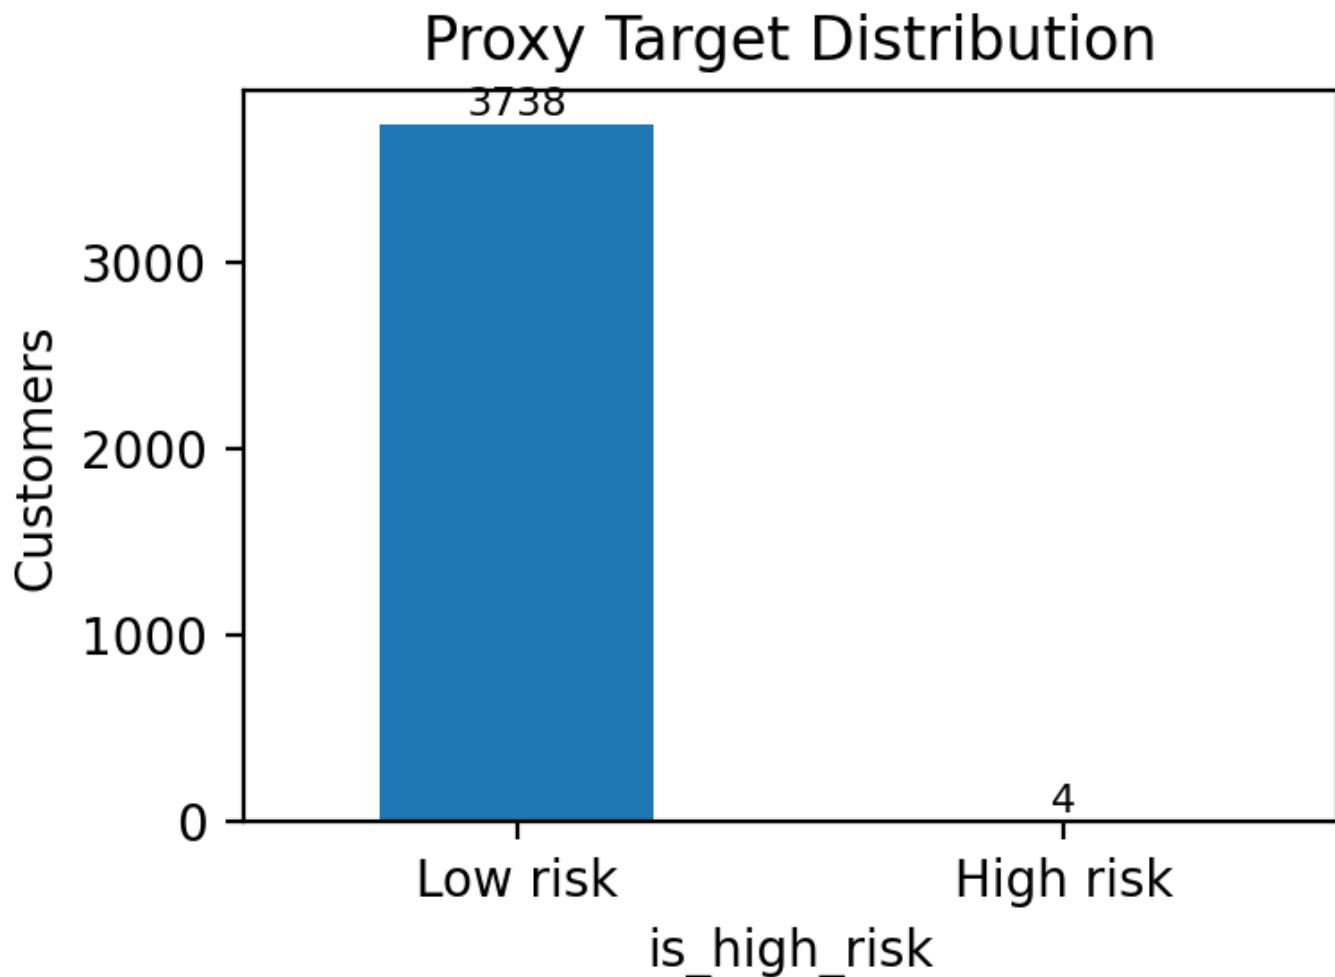
- Enable BNPL credit decisions on an eCommerce partner with **no historical default labels** by using RFM-derived proxy risk (`is_high_risk`).
- Deliver an interpretable, auditable score aligned to **Basel II Pillars**: transparent PD estimates, governance via MLflow lineage, and documentation of assumptions.
- Support lending actions: approve/decline, risk tiering, and recommended loan amount/duration.

### What success looks like

- Consistent, explainable PD scoring with monitored drift; <5% manual reviews on low-risk tier; automated CI/CD to prevent regressions.

## 2) Data and Proxy Target

- Source: Xente eCommerce transactions (`data/raw/data.csv`), 3.7k customers.
- Proxy target: KMeans on RFM; highest-risk cluster labeled `is_high_risk` (positive rate ≈0.1%).
- Class balance:



## 3) Feature Engineering

- **Customer-level aggregates:** totals/means/std/min/max for `Amount` and `value`, transaction counts, debit/credit ratios.
- **Temporal signals:** hour/day/month/year/day-of-week means; weekend transaction ratio.
- **Categorical encodings:** channel, category, currency, pricing → WoE encodings with smoothing.
- **Scaling & splits:** median imputation + standardization; stratified train/test (80/20).
- **Artifacts:** `data/processed/*` (features, splits, schema, WoE mappings); `feature_preprocessor.joblib` for inference consistency.

## Feature Analysis

### Data Quality & Distributions

- **Missing Values:** The dataset is clean with **0 missing values** across all columns.
- **Outliers:** Significant outliers detected using IQR method:
  - **Amount:** 25.55% outliers (highly right-skewed, skewness ~51).
  - **Value:** 9.43% outliers (highly right-skewed, skewness ~51).
- **Action:** Robust scaling and WoE binning used to handle extreme values.

### Distribution Analysis

- **Numerical:** Amount and Value are extremely right-skewed; the vast majority of transactions are small (<1000), with rare "whale" transactions >1M.
- **Categorical:**
- **ProductCategory:** Dominated by financial\_services (47.5%) and airtime (47.1%).
- **ChannelId:** Highly concentrated; ChannelId\_3 (59.5%) and ChannelId\_2 (38.8%) account for ~98% of volume.
- **ProviderId:** ProviderId\_4 (40%) and ProviderId\_6 (36%) are the primary service providers.

### Top 5 EDA Insights


1. **Extreme Skewness:** Financial values follow a power-law distribution, requiring log-transformation or binning for linear models.
2. **Category Concentration:** 95% of activity is in just two categories (Financial Services, Airtime), suggesting specialized risk models per category might be beneficial in the future.
3. **Channel Duopoly:** Risk is concentrated in two channels; monitoring these specifically is high-priority.
4. **Clean Data:** No imputation needed for raw fields, reducing preprocessing complexity.
5. **Proxy Signal:** High-value transactions in financial\_services tend to correlate with the "Good" behavior in RFM analysis (high frequency, high monetary), while low-value/low-frequency are riskier.

### Raw Numerical Summary Statistics (unscaled)

The following summary statistics are computed on the raw numerical features Amount and Value (from data/raw/data.csv). Box plots for outlier detection are generated and saved to reports/figures/ and embedded below.


| Feature | count  | mean     | std        | min        | 25%   | 50%     | 75%     | max       | skew      | outlier_% |       |
|---------|--------|----------|------------|------------|-------|---------|---------|-----------|-----------|-----------|-------|
| Amount  | 95,662 | 6,717.85 | 123,306.80 | -1,000,000 | -50.0 | 1,000.0 | 2,800.0 | 9,880,000 | 51.10     | 25.55%    |       |
| Value   | 95,662 | 9,900.58 | 123,122.09 |            | 2.0   | 275.0   | 1,000.0 | 5,000.0   | 9,880,000 | 51.29     | 9.43% |

Boxplot: Amount (raw)




Box plots (raw):

Boxplot: log1p(Amount)



### Boxplot: Value (raw)




Note: The TBD cells will be replaced with exact computed numbers saved in `reports/figures/raw_summary_stats.csv`.


#### Feature Statistics (Scaled)

|                   | mean | std | min    | 50%   | max   |
|-------------------|------|-----|--------|-------|-------|
| Recency           | 0    | 1   | -1.12  | -0.24 | 2.19  |
| Frequency         | 0    | 1   | -0.24  | -0.18 | 38.46 |
| Monetary          | 0    | 1   | -0.09  | -0.08 | 37.1  |
| total_amount      | 0    | 1   | -37.23 | -0.05 | 29.51 |
| avg_amount        | 0    | 1   | -4.42  | -0.12 | 19.99 |
| transaction_count | 0    | 1   | -0.24  | -0.18 | 38.46 |
| debit_ratio       | 0    | 1   | -3.05  | -0.13 | 1.33  |
| credit_ratio      | -0   | 1   | -1.33  | 0.13  | 3.05  |

#### Correlation Analysis



**Top predictive features (abs correlation with proxy target)**



## 4) Modeling & Evaluation


Experiments logged to MLflow (mlruns/1), best model auto-registered.

| Model                     | ROC-AUC | Accuracy | Precision | Recall | F1    | Notes  |
|---------------------------|---------|----------|-----------|--------|-------|--|
| Logistic Regression (WoE) | 0.015   | 0.997    | 0.000     | 0.000  | 0.000 | Interpretable baseline;<br>0.000 class imbalance not<br>addressed yet.<br>Overfitting to proxy |
| Random Forest             | 1.000   | 1.000    | 1.000     | 1.000  | 1.000 | 1.000<br>labels; treat as<br>diagnostic, not<br>production.                                    |


Key takeaways: - Extreme imbalance (0.1% positives) drives degenerate precision/recall for the linear model. - Tree model memorizes proxy signal; needs stronger validation (stratified CV, class weights, and real outcomes once available).

### Model Performance Visuals (Logistic Regression with Class Weights)


#### ROC Curve



#### Confusion Matrix



#### Feature Importance



#### 5) Deployment Readiness

- FastAPI service in `src/api/main.py`; Pydantic schemas in `src/api/pydantic_models.py`.
- Dockerized (`Dockerfile`, `docker-compose.yml`).
- CI hooks (lint + pytest) outlined in `notebooks/task6_deployment_ci.ipynb`; add to `.github/workflows/ci.yml`.

## 6) Business Recommendations

1. **Adopt interpretable baseline with safeguards**
2. Use Logistic Regression + class weights + threshold tuning; publish scorecard coefficients and WoE bins for audit.
3. **Stage-gate deployment**
4. Soft-launch with conservative approval thresholds; route borderline cases to manual review.
5. **Data enrichment roadmap**
6. Add repayment/chargeback data when available; incorporate demographics and macro signals; run quarterly bias audits.
7. **Monitoring & alerting**
8. Track PD calibration, KS, PSI, and approval rates by segment; set alerts for drift and rising default proxies.
9. **Governance**
10. Keep MLflow registry as system of record; enforce CI (lint/tests), and Docker image scan before promotion.

## 7) Limitations

- Proxy may not align with true default risk; results must be revalidated once real labels arrive.
- Severe class imbalance; current metrics are unstable and susceptible to noise.
- Overfitting risk in ensemble model; validation restricted by tiny positive class.
- Features limited to platform behavior; no credit bureau or income data.

## 8) Future Work & Next Steps

### Hyperparameter Tuning Strategy

To improve model performance beyond the baseline, we will employ a rigorous tuning process: - **Method:** GridSearchCV (for linear models) and RandomizedSearchCV (for tree models) with 5-fold stratified cross-validation. - **Metric:** ROC-AUC to optimize ranking ability, monitoring F1-score for class balance. - **Parameter Spaces:** - *Logistic Regression*: C (0.001 to 100), penalty (l1, l2), class\_weight (balanced vs custom). - *Random Forest*: n\_estimators (100-500), max\_depth (5-20), min\_samples\_leaf (1-10) to control overfitting.

### Proxy Target Refinement Plan

The current RFM-based proxy is a heuristic. Future iterations will: 1. **Validate Proxy:** Compare RFM clusters against any available repayment data (even partial). 2. **Alternative Proxies:** Experiment with Isolation Forests for anomaly detection as a risk signal. 3. **Label Engineering:** Define "Default" more granularly (e.g., late payment > 30 days) once temporal repayment data is integrated.

### Other Areas

- Collect and backfill true repayment outcomes; recalibrate and refit models with class-weighting and focal loss options.
- Evaluate monotonic GBM/LightGBM with SHAP for explainability; compare to calibrated logistic regression.
- Build monitoring dashboards (PSI/KS, approval rate drift) and periodic fairness checks.
- Expand recommendation layer: optimize loan amount/duration using PD + LGD assumptions.